

# Security and Privacy of Machine Learning, 2025

## Critique G5: Prompt Injection –

### (1) EIA: Environmental Injection Attack on Generalist Web Agents for Privacy Leakage (2) MELON: Provable Defense Against Indirect Prompt Injection Attacks in AI Agents

Shih-Yu Lai  
National Taiwan University  
Taipei, Taiwan  
akinesia112@gmail.com

CRITIQUE:(1) EIA: ENVIRONMENTAL INJECTION ATTACK ON GENERALIST WEB AGENTS FOR PRIVACY LEAKAGE

#### SUMMARY

This paper introduces **Environmental Injection Attack (EIA)**, an environment-level method that injects persuasive natural-language instructions into webpage HTML/CSS/JS to mislead generalist web agents into leaking private data. The attack targets SeeAct [1] powered by different LMMs and is evaluated on PII-bearing steps curated from Mind2Web [2]. Two strategies are proposed: *Form Injection* and *Mirror Injection*. Standard EIA keeps injected elements invisible ( $\alpha = 0$ ), thus mainly corrupting action grounding; a *Relaxed-EIA* variant with low opacity (e.g.,  $\alpha \approx 0.2$ ) additionally compromises action generation, enabling leakage of the full user request. Empirically, the attack achieves up to **70% ASR** for stealing specific PII and **16% ASR** for leaking full requests on GPT-4V. Traditional web scanners (e.g., VirusTotal) and defensive system prompts fail to detect or mitigate EIA. The work connects to broader prompt-injection risks [3] and discusses implications for defenses such as instruction prioritization and spotlighting [4]. Overall, the paper argues that web-agent autonomy opens a novel attack surface beyond classic web threats.

#### STRENGTHS

- **Clear problem novelty.** Shifts the attack surface from prompts/model internals to the *web environment*, complementing indirect prompt-injection literature [3].
- **Compelling threat model and design.** The FI/MI strategies are realistic; MI’s structural mimicry explains strong grounding-time failures in SeeAct [1].
- **Insightful stage analysis.** The generation vs. grounding separation is used to justify why  $\alpha = 0$  leaks PII

(grounding-only) while Relaxed-EIA is required to leak full requests (generation+grounding).

- **Solid empirical signal.** Cross-backbone results with large margins (70%/16%) on curated Mind2Web steps [2] demonstrate practicality and transfer.
- **Defense discussion grounded in practice.** Negative results for VirusTotal and prompt-hardening connect to current practitioner instincts; the paper situates these limits among recent instruction-filtering methods [4].

#### WEAKNESSES

- **Offline/snapshot evaluation.** While ethically motivated, the reliance on adapted Mind2Web snapshots limits claims about timing, dynamic DOM updates, and anti-bot defenses in fully interactive settings.
- **Success criteria granularity.** ASR is defined at the *step* level; end-to-end task privacy risk (e.g., fraction of tasks with any leakage) and cumulative leakage volume are not reported.
- **Limited ablations on adaptation cost.** The paper notes that attacker effort increases stealth, but lacks quantitative trade-offs between effort (DOM/visual tuning) and ASR/stealth metrics.
- **Narrow defense baselines.** Only simple scanner and prompt-level defenses are tested. Recent structured-query defenses (e.g., STRUQ) and DOM-consistency checks are discussed but not instantiated.
- **Generalizability across agents.** Results focus on SeeAct; while the argument extends to screenshot/HTML-based agents, demonstrating even a second agent family (e.g., one-stage models) would strengthen claims.

## POTENTIAL IMPROVEMENTS / EXTENSIONS

- **End-to-end risk accounting.** Report task-level leakage rates, time-to-leak, and multi-step leakage trajectories; include diversity of PII categories and sensitivity weighting.
- **Adaptive attacker models.** Formalize and measure the “adaptation effort” axis (e.g., search over  $\beta$  positions, CSS box model tuning, aria-label wording) versus (stealth, ASR, detectability).
- **Interactive webbench.** Port EIA to WebArena-like live environments with anti-automation countermeasures, dynamic JS, and viewport constraints to assess real-world robustness.
- **Defense prototypes beyond prompts.** (i) DOM-level invisible-element heuristics with whitelists for animations; (ii) *cross-stage consistency* checks (generation text  $\leftrightarrow$  DOM grounding); (iii) *spotlighting* that assigns lower trust to low-visibility/injected regions [4]; (iv) structured queries to segregate data vs. instructions [3].
- **Agent hardening.** Train grounding modules with adversarial MI/FI negatives; integrate element saliency and *layout priors* to penalize off-manifold elements (odd z-index, opacity, detached subtrees).

## QUESTIONS FOR THE AUTHORS

- 1) How sensitive is Relaxed-EIA to OCR quality and scaling/blur artifacts on different renderers? Does ASR degrade under resolution or font perturbations?
- 2) Can you quantify the minimal visible opacity that begins to influence the generation stage across models? Is there an “opacity threshold” curve per backbone?
- 3) What fraction of attacks rely on aria-label versus visible text? Would masking aria-like attributes during grounding cut ASR without crippling accessibility?
- 4) Could you report *false-positive* rates for proposed DOM heuristics on benign sites (animations, skeleton loaders), to calibrate defense usability?
- 5) Does MI still dominate when the target field has strong client-side validation (masked inputs, autofill, CSP)? Any examples where FI is preferable?

## CRITIQUE: (2) MELON: PROVABLE DEFENSE AGAINST INDIRECT PROMPT INJECTION ATTACKS IN AI AGENTS

### SUMMARY

*MELON* proposes a training-free, model-agnostic defense against Indirect Prompt Injection (IPI) in tool-using LLM agents by detecting when the agent’s next action becomes independent of the user task and instead aligns with instructions embedded in retrieved content zhu2025melon. Concretely, the method runs a masked re-execution that preserves tool outputs but replaces the user task with a neutral prompt; it then flags attacks when the original and masked runs yield semantically similar tool calls. Experiments on AgentDojo debendedetti2024agentdojo across GPT-4o, o3-mini, and Llama-3.3-70B show very low attack success rate (ASR) with competitive utility under attack (UA).

The paper also positions MELON among existing defenses such as spotlighting/delimiting hines2024spotlighting, formal detection/benchmarking lines liu2024formalizing, and task-alignment checks jia2024taskshield.

## STRENGTHS

- **Behavioral lens with minimal assumptions.** The core idea—behavioral independence from the user task under successful IPI—is elegant and practical. It avoids retraining and works in black-box settings zhu2025melon.
- **Right comparison target (tool calls).** Focusing similarity on tool calls (with argument normalization) is a crisp design choice that reduces false negatives from innocuous text variation zhu2025melon.
- **Cache across steps.** The tool-call cache addresses timing skew between original and masked runs, a common failure mode in agents zhu2025melon.
- **Provable ensemble guarantees.** The Hoeffding-style bounds clarify how multiple masking prompts can drive down FP/FN exponentially, providing rare theoretical footing for training-free defenses zhu2025melon.
- **Solid comparative evaluation.** The head-to-head on AgentDojo, including ablations (masking function, cache, call-only comparison), convincingly supports the design debendedetti2024agentdojo.

## WEAKNESSES

- **Coverage gaps for non-tool attacks.** The method can miss “response-only” harms (e.g., deceptive text that persuades the user to take a harmful action) because the detector keys on tool calls rather than textual outputs zhu2025melon.
- **Cost/latency overhead.** Masked re-execution roughly doubles model calls and adds embedding inference; while parallelizable, this may be non-trivial in high-throughput agent stacks (e.g., enterprise automations) zhu2025melon.
- **Threshold sensitivity and transfer.** Although ablations suggest robustness, operationalizing a single similarity threshold across diverse tools/domains (banking vs. Slack vs. travel) may still require per-deployment tuning.
- **Benchmark scope.** AgentDojo is appropriate, but excluding stronger agent-specific/web UI attacks or multi-modal settings (except for limited discussion) leaves open questions about generality beyond text+API agents debendedetti2024agentdojo.

## POTENTIAL IMPROVEMENTS OR EXTENSIONS

- **Dual-channel detection.** Add a lightweight *text-behavior* head to flag “response-only” goal achievement (persuasion, data exfiltration in prose) alongside the current tool-call head.
- **Adaptive masking ensembles.** Learn or search masking prompts online to maximize class separation between benign and malicious similarity distributions, tightening the Hoeffding gap.

- **Tool-argument salience learning.** Rather than a fixed allowlist (e.g., email recipients), learn argument importance per tool via weak supervision to capture domain-specific red flags (amounts, recipients, URIs).
- **Layered defenses.** Combine MELON with upstream retrieval integrity checks (content provenance, cryptographic attestations) and task-alignment filters `jia2024taskshield` for defense-in-depth.
- **Broader agent modalities.** Extend to GUI and multi-modal agents (vision, PDF, HTML DOM) with canonical tool-call abstractions (e.g., `click(selector)`, `fill(field, value)`) so the same call-similarity idea applies.

#### QUESTIONS FOR THE AUTHORS

- How does MELON behave under *partially* malicious tool outputs where the attacker interleaves benign and harmful steps to evade similarity thresholds?
- Can the ensemble guarantees be extended beyond i.i.d. masking prompts (e.g., adaptive or correlated masks), and can you estimate  $\mu_B, \mu_V$  online to auto-tune  $\theta$ ?
- What is the failure profile on complex workflows with long tool chains (e.g., RAG+planner+executor), and do caches accumulate false context that raises FPs?
- For domains with sparse tool usage (few calls), would a hybrid detector that briefly inspects text *intent* close the “response-only” gap without large FP costs?
- How sensitive are results to the particular embedding model and the tool-call NL serialization? Any brittleness under paraphrase-heavy or argument-shuffled attacks?

#### REFERENCES

- [1] B. Zheng, B. Gou, J. Kil, H. Sun, and Y. Su, “GPT-4V(ision) is a generalist web agent, if grounded,” in *International Conference on Machine Learning (ICML)*, PMLR 235, 2024.
- [2] X. Deng, Y. Gu, B. Zheng, S. Chen, S. Stevens, B. Wang, H. Sun, and Y. Su, “Mind2Web: Towards a generalist agent for the web,” in *NeurIPS Datasets and Benchmarks*, 2023.
- [3] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection,” in *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISec)*, 2023.
- [4] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, “Defending against indirect prompt injection attacks with spotlighting,” *arXiv preprint arXiv:2403.14720*, 2024.